

Image Caption Generation

201368087

April 27, 2020

Contents

1	Text Preparation	1
2	Network Outputs	2
2.1	LSTM	2
2.2	RNN	4
3	Comparison of RNN vs LSTM	6
4	Additional Text Annotation Files	7

1 Text Preparation

One of the consequences of lemmatizing the words, means that any suffixes are removed from the words. This allows the model to focus on recognising objects within the image, and describing the object. Whereas without lemmatization, the model will have to also take into the account complex linguistical properties such as grammar, which is hard to model efficiently. Having to do two complex functions within one set of model and weights can be tricky, lemmatization can help to overcome that. Requiring the model to learn both grammar, and describing the image might go against abstraction principle; why should the model be required to do both tasks? It would make more sense for the model to just learn how to caption, and another program, possibly another model, could fix any grammatical mistakes. Especially since lemmatization is a very complex task, as it's not as simple as adding a suffix; in the case of words such as 'good' there is no 'gooder', as one example. Also, lemmatization keeps the vocabulary short, meaning that the model doesn't require to learn how each word can be used, and in which context. Mainly, the advantage of lemmatization is that it focusses the model's attention on relationship between words.

However, there are certain issues with lemmatization, as the generated captions will make often no grammatical sense, thus making it harder to evaluate, as there is no simple way to make the sentence be grammatically correct. However, leaving the sentence with grammatical errors, whilst achieving the goal of being close to the original captions, loses in terms of building coherent sentences a human would make. This makes it hard to say whether the model was successful or not. Also, at times lemmatization does not always lead to better performance, therefore it is best to not use it, unless the model's performance cannot be improved, thus trying lemmatization could help the performance. Furthermore, some sentiment analysis tools differentiate between different words based on the context and give a rating based on the context, thus lemmatization would be harmful if it were to be used in conjunction with tools like VADER.

To conclude, lemmatization could remove some useful words, or it could remove a lot of clutter thus leading to a better model. It is hard to say what impact lemmatization would have on the model, without trying it.

2 Network Outputs

2.1 LSTM



Reference Captions:

a child playing on the monkey bars at a playground with an adult <unk>
a kid hangs upside down on the monkey bars at a playground
a young children hanging upside at the playground
girl hanging upside down on blue monkey bars
a girl hangs on the bars

Generated Caption:

cross-country choppy other courtyard offering tiger treat martial guide awards melting foliage elder backward elder rottweiler sits mexican college stoop
BLUE score: 0.000000



Reference Captions:

collie jumping over a training hurdle that is on the grass
a brown and white sheltie leaping over a rail
a collie is running through an obstacle course
collie making a jump over a yellow hurdle
the dog is leaping over a hurdle

Generated Caption:

military foliage conversing conversing directions threw somersault also reaches boulder museum climb art jackets kicked ocean wristbands lit dancer costumes
BLUE score: 0.000000

Figure 1: Initial Caption



Reference Captions:

a child playing on the monkey bars at a playground with an adult <unk>
a kid hangs upside down on the monkey bars at a playground
a young children hanging upside at the playground
girl hanging upside down on blue monkey bars
a girl hangs on the bars

Generated Caption:

a man in a red shirt is standing on a bench
BLUE score: 0.476851



Reference Captions:

collie jumping over a training hurdle that is on the grass
a brown and white sheltie leaping over a rail
a collie is running through an obstacle course
collie making a jump over a yellow hurdle
the dog is leaping over a hurdle

/usr/local/lib/python3.6/dist-packages/nltk/translate/bleu_score.py:490: UserWarning:
Corpus/Sentence contains 0 counts of 4-gram overlaps.
BLEU scores might be undesirable; use SmoothingFunction().
warnings.warn(_msg)

Generated Caption:

a dog is running through the grass
BLUE score: 0.709793

Figure 2: After 1 Epoch



Reference Captions:
 a child playing on the monkey bars at a playground with an adult <unk>
 a kid hangs upside down on the monkey bars at a playground
 a young children hanging upside at the playground
 girl hanging upside down on blue monkey bars
 a girl hangs on the bars

Generated Caption:
 a boy in a blue shirt and blue jeans is jumping into a pool
 BLUE score: 0.462910



Reference Captions:
 collie jumping over a training hurdle that is on the grass
 a brown and white sheltie leaping over a rail
 a collie is running through an obstacle course
 collie making a jump over a yellow hurdle
 the dog is leaping over a hurdle

Generated Caption:
 a dog is running through a field
 BLUE score: 0.611538

Figure 3: After 2 Epoch



Reference Captions:
 a child playing on the monkey bars at a playground with an adult <unk>
 a kid hangs upside down on the monkey bars at a playground
 a young children hanging upside at the playground
 girl hanging upside down on blue monkey bars
 a girl hangs on the bars

Generated Caption:
 a young boy in a blue shirt is jumping into a pool
 BLUE score: 0.317023



Reference Captions:
 collie jumping over a training hurdle that is on the grass
 a brown and white sheltie leaping over a rail
 a collie is running through an obstacle course
 collie making a jump over a yellow hurdle
 the dog is leaping over a hurdle

Generated Caption:
 a dog jumps over a hurdle
 BLUE score: 0.499152

Figure 4: After 3 Epoch



Reference Captions:
 a child playing on the monkey bars at a playground with an adult <unk>
 a kid hangs upside down on the monkey bars at a playground
 a young children hanging upside at the playground
 girl hanging upside down on blue monkey bars
 a girl hangs on the bars

Generated Caption:
 a girl in a pink shirt is jumping on a trampoline
 BLUE score: 0.309637



Reference Captions:
 collie jumping over a training hurdle that is on the grass
 a brown and white sheltie leaping over a rail
 a collie is running through an obstacle course
 collie making a jump over a yellow hurdle
 the dog is leaping over a hurdle

Generated Caption:
 a dog jumps over a hurdle
 BLUE score: 0.499152

Figure 5: After 4 Epoch



Reference Captions:
 a child playing on the monkey bars at a playground with an adult <unk>
 a kid hangs upside down on the monkey bars at a playground
 a young children hanging upside at the playground
 girl hanging upside down on blue monkey bars
 a girl hangs on the bars

Generated Caption:
 a boy in a blue shirt is jumping on a trampoline
 BLUE score: 0.550620



Reference Captions:
 collie jumping over a training hurdle that is on the grass
 a brown and white sheltie leaping over a rail
 a collie is running through an obstacle course
 collie making a jump over a yellow hurdle
 the dog is leaping over a hurdle

Generated Caption:
 a dog jumps over a hurdle
 BLUE score: 0.499152

Figure 6: After 5 Epoch

2.2 RNN



Reference Captions:
 a child playing on the monkey bars at a playground with an adult <unk>
 a kid hangs upside down on the monkey bars at a playground
 a young children hanging upside at the playground
 girl hanging upside down on blue monkey bars
 a girl hangs on the bars

Generated Caption:
 crowd patches jogger fenced dunking guys skirt goal flag gated violin breath agility piles bearded snow-covered pull world group listening
 BLUE score: 0.000000



Reference Captions:
 collie jumping over a training hurdle that is on the grass
 a brown and white sheltie leaping over a rail
 a collie is running through an obstacle course
 collie making a jump over a yellow hurdle
 the dog is leaping over a hurdle

Generated Caption:
 crowd patches jogger fenced dunking guys skirt goal flag gated violin breath agility piles bearded snow-covered pull world group listening
 BLUE score: 0.000000

Figure 7: Initial Caption



Reference Captions:
 a child playing on the monkey bars at a playground with an adult <unk>
 a kid hangs upside down on the monkey bars at a playground
 a young children hanging upside at the playground
 girl hanging upside down on blue monkey bars
 a girl hangs on the bars

Generated Caption:
 a girl in a pink shirt is jumping off a rock
 BLUE score: 0.268153



Reference Captions:
 collie jumping over a training hurdle that is on the grass
 a brown and white sheltie leaping over a rail
 a collie is running through an obstacle course
 collie making a jump over a yellow hurdle
 the dog is leaping over a hurdle

Generated Caption:
 a dog runs through the water
 BLUE score: 0.691149

Figure 8: After 1 Epoch



Reference Captions:
 a child playing on the monkey bars at a playground with an adult <unk>
 a kid hangs upside down on the monkey bars at a playground
 a young children hanging upside at the playground
 girl hanging upside down on blue monkey bars
 a girl hangs on the bars

Generated Caption:
 a boy in a blue shirt is riding a dirt bike
 BLUE score: 0.476851



Reference Captions:
 collie jumping over a training hurdle that is on the grass
 a brown and white sheltie leaping over a rail
 a collie is running through an obstacle course
 collie making a jump over a yellow hurdle
 the dog is leaping over a hurdle

Generated Caption:
 a dog jumps over a hurdle
 BLUE score: 0.499152

Figure 9: After 2 Epoch



Reference Captions:
 a child playing on the monkey bars at a playground with an adult <unk>
 a kid hangs upside down on the monkey bars at a playground
 a young children hanging upside at the playground
 girl hanging upside down on blue monkey bars
 a girl hangs on the bars

Generated Caption:
 a little girl in a red shirt is jumping off a swing
 BLUE score: 0.500000



Reference Captions:
 collie jumping over a training hurdle that is on the grass
 a brown and white sheltie leaping over a rail
 a collie is running through an obstacle course
 collie making a jump over a yellow hurdle
 the dog is leaping over a hurdle

Generated Caption:
 a dog jumps over a hurdle
 BLUE score: 0.499152

Figure 10: After 3 Epoch



Reference Captions:
 a child playing on the monkey bars at a playground with an adult <unk>
 a kid hangs upside down on the monkey bars at a playground
 a young children hanging upside at the playground
 girl hanging upside down on blue monkey bars
 a girl hangs on the bars

Generated Caption:
 a child in a blue shirt is jumping off a swing
 BLUE score: 0.309637



Reference Captions:
 collie jumping over a training hurdle that is on the grass
 a brown and white sheltie leaping over a rail
 a collie is running through an obstacle course
 collie making a jump over a yellow hurdle
 the dog is leaping over a hurdle

Generated Caption:
 a dog jumps over a hurdle
 BLUE score: 0.499152

Figure 11: After 4 Epoch



Reference Captions:
 a child playing on the monkey bars at a playground with an adult <unk>
 a kid hangs upside down on the monkey bars at a playground
 a young children hanging upside at the playground
 girl hanging upside down on blue monkey bars
 a girl hangs on the bars

Generated Caption:
 a boy in a blue shirt is swinging on a swing
 BLUE score: 0.550620



Reference Captions:
 collie jumping over a training hurdle that is on the grass
 a brown and white sheltie leaping over a rail
 a collie is running through an obstacle course
 collie making a jump over a yellow hurdle
 the dog is leaping over a hurdle

Generated Caption:
 a dog jumps over a hurdle
 BLUE score: 0.499152

Figure 12: After 5 Epoch

3 Comparison of RNN vs LSTM



Reference Captions:
 a person in a red shirt climbing up a rock face covered in <unk> handles
 a man in a pink shirt climbs a rock face
 a man is rock climbing high in the air
 a rock climber practices on a rock climbing wall
 a rock climber in a red shirt

Generated Caption:
 a man is climbing a rock wall
 BLUE score: 0.660537

(a) LSTM



Reference Captions:
 a person in a red shirt climbing up a rock face covered in <unk> handles
 a man in a pink shirt climbs a rock face
 a man is rock climbing high in the air
 a rock climber practices on a rock climbing wall
 a rock climber in a red shirt

Generated Caption:
 a man in a blue shirt climbs a rock wall
 BLUE score: 0.681612

(b) RNN

Figure 13: Caption of a rock climbing image

The loss for both the RNN and LSTM was between 2.9 and 2.1, and seemed to oscillate around that range, with the RNN ending at 2.5, and LSTM at 2.2. At first, both of the models obviously started high at around 8, which at the end of the first epoch went down to 3.0. After that, it seemed to go up and down at around 2.9-2.1 for both models. However, the

RNN would have more variance within epochs. For example, the highest loss for epoch 5 of RNN was 2.5363 whilst lowest one was 2.1525, which is a large difference in this context. This is further magnified by the fact that the LSTM model's fifth epoch ranged between 2.1670 and 2.3768. This trend can be seen in other epochs but slightly less extreme.

Considering everything, the loss between the two models is very close, and you would be nitpicking to find any major differences. What the loss boils down to is both the models were consistently improving the loss for the first epoch, and afterwards the loss stagnated and oscillated between a range of values, rather than uniformly decreasing. Conclusively, both of the models struggle in calculating the loss, which is to be expected given that evaluating captions is no easy task.

Currently, a more pragmatic approach would be beneficial by looking at the quality of the captions in a less algorithmic way. Personally, a lot of the captions have grammatical structure that would pass by a sentence made by a human. A lot of the captions also accurately depict what is displayed in the image. As an example, the dog image was captioned pretty much spot on. The only issue with it, is the lack of detail. The reference captions often include more detail, such as a breed of a dog, shirt colour of a person, and so on. Both of the models seem opposed to adding in detail.

Furthermore, given the images in the RNN gave a more descriptive caption, however an erroneous one as the person in question does not have a blue shirt. On the other hand, the LSTM managed to give a 100% correct caption, but much shorter with less detail. However, the BLEU score favours the inaccurate, yet more advanced caption given by the RNN. This shows that the BLEU scores are not perfect, as, personally, the less descriptive but accurate statement should be scored way higher given that it is completely plausible that a human would give such a caption.

In general, I have found out that the BLEU scores can be quite misleading in certain cases. This is due to the fact that the BLUE scores strictly calculates how close it is to the reference captions, thus leading to some confusing scores. For example, for the dog image the caption 'a dog jumps over a hurdle' gets a lower score than 'a dog runs through water'. This is partly due to the weights I have chosen for the n-grams. I have decided to use $weights = (0.5, 0.25, 0.15, 0.1)$, as I did not want to put too much emphasis on the 3/4-gram as in my testing, 'a dog jumps over a hurdle' would get an even lower score. The cause of that would be that the model doesn't seem to favour descriptive captions. This is reinforced by the fact that in both the RNN and LSTM flavours, the 'a dog jumps over the hurdle' caption doesn't change for 3 epochs.

4 Additional Text Annotation Files

There are two files; namely ExpertAnnotations.txt and CrowdFlowerAnnotations.txt. ExpertAnnotations.txt contains a table of images and captions from other images, which the expert ranked from 1-4 as to whether a caption from that different image is similar to the image in question. In similar vein, the CrowdFlowerAnnotations.txt file contains the number of yeses and number of noes whether a caption from a different image describes another image. A percentage of yesses is also recorded.

These files have similar functions, which solve a hard problem when it comes to models in this domain, which is encapsulating the wide variety of ways to describe a given image. There is no one fixed solution to the problem, and there are countless ways to describe an image that are perfectly fine, or at least convey the basic message. Even if the model manages to successfully come close to the captions used in training, there are other ways that describe the image just as well, but with different level of detail or tone etc. Thus, the hardest part would be to get a large sample size to effectively narrow down a way of describing that object. Trying to encapsulate all the possible ways gets even harder when you consider words that portray different emotions like 'badly' or 'awesome'. Furthermore, the amount of detail given may vary from person to person, as seen previously with *figure 13*. Such variance is proven to be tricky to evaluate, thus to improve confidence of the model, it might be more beneficial to define what not to do, rather than what to do.

These files would help with this issue, as they could be used as a similarity ranking between the different forms of captions. The files can also be used in a unique way comparatively to the captions attached to the images, as they can show which captions they shouldn't be close to. For example, if a caption is close in nature to a caption with a lot of 'No's given by the experts/crowd then that can be considered an inadequate caption and possibly count it in the loss function. This also gives a better way to evaluate the model, as if the caption is somewhat close to the attached captions, but close to captions it shouldn't be close to, then it might suggest that the model still hasn't fully recognised the object in question, and a lot of the score can be attributed to containing stop words like 'the'.

References